

Introduction to PyTorch

October 27, 2022

1 Lecture 16: PyTorch Crash Course

1.1 Tensor Initialization

See the full API at <https://pytorch.org/docs/stable/tensors.html>

```
[1]: import torch
import numpy as np
```

```
[2]: x=torch.tensor([1,2,3,4])
print(x)
```

```
tensor([1, 2, 3, 4])
```

```
[3]: x.shape
```

```
[3]: torch.Size([4])
```

```
[4]: type(x)
```

```
[4]: torch.Tensor
```

```
[5]: x.dtype
```

```
[5]: torch.int64
```

```
[6]: x=torch.tensor([1,2,3,4]).type(torch.float)
x
```

```
[6]: tensor([1., 2., 3., 4.])
```

```
[7]: x.dtype
```

```
[7]: torch.float32
```

```
[8]: x=torch.FloatTensor([1,2,3,4])
x
```

```
[8]: tensor([1., 2., 3., 4.])
```

```
[9]: x.dtype
```

```
[9]: torch.float32
```

```
[10]: y = torch.FloatTensor([[1,2,3,4],[5,6,7,8]])  
y
```

```
[10]: tensor([[1., 2., 3., 4.],  
          [5., 6., 7., 8.]])
```

```
[11]: torch.zeros(5)
```

```
[11]: tensor([0., 0., 0., 0., 0.])
```

```
[12]: torch.ones(5)
```

```
[12]: tensor([1., 1., 1., 1., 1.])
```

```
[13]: torch.rand(5)
```

```
[13]: tensor([0.2323, 0.2190, 0.2919, 0.2940, 0.9902])
```

```
[14]: torch.eye(5)
```

```
[14]: tensor([[1., 0., 0., 0., 0.],  
          [0., 1., 0., 0., 0.],  
          [0., 0., 1., 0., 0.],  
          [0., 0., 0., 1., 0.],  
          [0., 0., 0., 0., 1.]])
```

```
[15]: torch.FloatTensor(np.random.randn(5))
```

```
[15]: tensor([-0.9962, -0.0668,  0.1160,  0.5321,  2.0040])
```

1.2 Linear Algebra

```
[16]: x=torch.FloatTensor([[1,2,3,4]])  
z=torch.FloatTensor([[1,2,3,4],[5,6,7,8]])  
print("x=",x)  
print("z=",z)
```

```
x= tensor([[1., 2., 3., 4.]])  
z= tensor([[1., 2., 3., 4.],  
          [5., 6., 7., 8.]])
```

```
[17]: 5*x
```

```
[17]: tensor([[ 5., 10., 15., 20.]])
```

```
[18]: x+5
```

```
[18]: tensor([[6., 7., 8., 9.]])
```

```
[19]: x+x
```

```
[19]: tensor([[2., 4., 6., 8.]])
```

```
[20]: x.T
```

```
[20]: tensor([[1.],  
          [2.],  
          [3.],  
          [4.]])
```

```
[21]: x@x.T
```

```
[21]: tensor([[30.]])
```

```
[22]: x.T@x
```

```
[22]: tensor([[ 1.,  2.,  3.,  4.],  
          [ 2.,  4.,  6.,  8.],  
          [ 3.,  6.,  9., 12.],  
          [ 4.,  8., 12., 16.]])
```

```
[23]: x@z.T
```

```
[23]: tensor([[30., 70.]])
```

1.3 Broadcasting

```
[24]: x=torch.FloatTensor([[1,2,3,4]])  
      y=torch.FloatTensor([[-1,-2]]).T  
      z=torch.FloatTensor([[1,2,3,4],[5,6,7,8]])  
      print("x=",x)  
      print("y=",y)  
      print("z=",z)
```

```
x= tensor([[1., 2., 3., 4.]])  
y= tensor([[-1.],  
          [-2.]])  
z= tensor([[1., 2., 3., 4.],  
          [5., 6., 7., 8.]])
```

```
[25]: x+y
```

```
[25]: tensor([[ 0.,  1.,  2.,  3.],  
            [-1.,  0.,  1.,  2.]])
```

```
[26]: x*y
```

```
[26]: tensor([[ -1., -2., -3., -4.],  
            [-2., -4., -6., -8.]])
```

```
[27]: x+z
```

```
[27]: tensor([[ 2.,  4.,  6.,  8.],  
            [ 6.,  8., 10., 12.]])
```

```
[28]: x*z
```

```
[28]: tensor([[ 1.,  4.,  9., 16.],  
            [ 5., 12., 21., 32.]])
```

```
[29]: y+z
```

```
[29]: tensor([[0., 1., 2., 3.],  
            [3., 4., 5., 6.]])
```

```
[30]: y*z
```

```
[30]: tensor([[ -1., -2., -3., -4.],  
            [-10., -12., -14., -16.]])
```

1.4 Elementwise Functions

```
[31]: x=torch.FloatTensor([1,2,-3,-4])  
x
```

```
[31]: tensor([[ 1.,  2., -3., -4.]])
```

```
[32]: torch.pow(x,2)
```

```
[32]: tensor([[ 1.,  4.,  9., 16.]])
```

```
[33]: x**2
```

```
[33]: tensor([[ 1.,  4.,  9., 16.]])
```

```
[34]: torch.exp(x)
```

```
[34]: tensor([[2.7183, 7.3891, 0.0498, 0.0183]])
```

```
[35]: torch.log(x)
```

```
[35]: tensor([[0.0000, 0.6931,    nan,    nan]])
```

```
[36]: torch.sin(x)
```

```
[36]: tensor([[ 0.8415,  0.9093, -0.1411,  0.7568]])
```

```
[37]: torch.abs(x)
```

```
[37]: tensor([[1., 2., 3., 4.]])
```

1.5 Matrix Functions

```
[38]: z = torch.FloatTensor([[2,1],[1,3]])  
z
```

```
[38]: tensor([[2., 1.],  
           [1., 3.]])
```

```
[39]: torch.det(z)
```

```
[39]: tensor(5.)
```

```
[40]: torch.inverse(z)
```

```
[40]: tensor([[ 0.6000, -0.2000],  
           [-0.2000,  0.4000]])
```

```
[41]: torch.trace(z)
```

```
[41]: tensor(5.)
```

1.6 Aggregation Functions

```
[42]: x=torch.FloatTensor([1,2,3,4])  
z=torch.FloatTensor([1,2,3,4],[5,6,7,8])  
print("x=",x)  
print("z=",z)
```

```
x= tensor([1., 2., 3., 4.])  
z= tensor([1., 2., 3., 4.,  
          5., 6., 7., 8.])
```

```
[43]: x.sum()
```

```
[43]: tensor(10.)
```

```
[44]: z.sum()
```

```
[44]: tensor(36.)
```

```
[45]: z.sum(axis=0)
```

```
[45]: tensor([ 6.,  8., 10., 12.])
```

```
[46]: z.sum(axis=1)
```

```
[46]: tensor([10., 26.])
```

```
[47]: z.sum(axis=1,keepdims=True)
```

```
[47]: tensor([[10.],  
           [26.]])
```

```
[48]: x.mean()
```

```
[48]: tensor(2.5000)
```

```
[49]: x.prod()
```

```
[49]: tensor(24.)
```

```
[50]: x.min()
```

```
[50]: tensor(1.)
```

```
[51]: x.max()
```

```
[51]: tensor(4.)
```

```
[52]: z.min()
```

```
[52]: tensor(1.)
```

```
[53]: v,ind = z.min(axis=0)  
      print(v)  
      print(ind)
```

```
tensor([1., 2., 3., 4.])  
tensor([0, 0, 0, 0])
```

```
[54]: v,ind = z.min(axis=1)
      print(v)
      print(ind)
```

```
tensor([1., 5.])
tensor([0, 0])
```

1.7 Automatic Differentiation

```
[55]: x=torch.FloatTensor([[1,2,3,4],[-1,-2,-3,-4]])
      w=torch.tensor([[[-1.0],[-2.0],[2.0],[1.0]],requires_grad=True)
      b=torch.tensor(-1.0,requires_grad=True)
      y=torch.tensor([[1.0],[0.0]])
      print("x=",x)
      print("w=",w)
      print("b=",b)
```

```
x= tensor([[ 1.,  2.,  3.,  4.],
           [-1., -2., -3., -4.]])
w= tensor([[[-1.],
           [-2.],
           [ 2.],
           [ 1.]], requires_grad=True)
b= tensor(-1., requires_grad=True)
```

```
[56]: x.requires_grad
```

```
[56]: False
```

```
[57]: w.requires_grad
```

```
[57]: True
```

```
[58]: b.requires_grad
```

```
[58]: True
```

```
[59]: a=x[0,:]*w +b
      a
```

```
[59]: tensor([4.], grad_fn=<AddBackward0>)
```

```
[60]: a.requires_grad
```

```
[60]: True
```

```
[61]: a.backward()
      print("grad_w:",w.grad,"grad_b:",b.grad)
```

```
grad_w: tensor([[1.],
               [2.],
               [3.],
               [4.]]) grad_b: tensor(1.)
```

```
[62]: a=x[0,:]*w +b
      a.backward()
      print("grad_w:",w.grad,"grad_b:",b.grad)
```

```
grad_w: tensor([[2.],
               [4.],
               [6.],
               [8.]]) grad_b: tensor(2.)
```

```
[63]: w.grad=0*w.grad
      b.grad=0*b
      a=x[0,:]*w +b
      a.backward()
      print("grad_w:",w.grad,"grad_b:",b.grad)
```

```
grad_w: tensor([[1.],
               [2.],
               [3.],
               [4.]]) grad_b: tensor(1., grad_fn=<MulBackward0>)
```

```
[64]: def mse(y,yhat):
      return torch.mean((y-yhat)**2)

      w.grad=0*w.grad
      b.grad=0*b

      yhat = x*w+b
      loss = mse(y,yhat)
      loss.backward()
      print("grad_w:",w.grad,"grad_b:",b.grad)
```

```
grad_w: tensor([[ 9.],
               [18.],
               [27.],
               [36.]]) grad_b: tensor(-3., grad_fn=<MulBackward0>)
```

1.8 Neural Network Modules

See the full API at <https://pytorch.org/docs/stable/nn.html>

```
[65]: X = torch.randn(1000,2)
      Y = -3*X[:,[0]] + 2*X[:,[1]]**2 + 3
      Y[:5,:]
```



```
[65]: tensor([[8.1985],
            [4.5534],
            [6.4657],
            [6.9933],
            [2.6479]])
```

```
[66]: import torch.nn as nn

class linear(nn.Module):
    def __init__(self,d,k):
        super(linear, self).__init__()
        self.w = nn.Parameter(torch.rand(d,k))
        self.b = nn.Parameter(torch.rand(1,k))

    def forward(self,x):
        return x@self.w + self.b

model = linear(2,1)
model.forward(X[:10,:])
```

```
[66]: tensor([[0.8807],
            [0.7144],
            [1.2655],
            [0.5497],
            [0.8091],
            [0.4771],
            [0.8583],
            [1.0352],
            [0.9309],
            [0.8387]], grad_fn=<AddBackward0>)
```

```
[67]: Yhat = model.forward(X)
loss = mse(Yhat,Y)
loss
```

```
[67]: tensor(35.4643, grad_fn=<MeanBackward0>)
```

```
[68]: class relu_mlp(nn.Module):
    def __init__(self,d,k):
        super(relu_mlp, self).__init__()
        self.l1 = nn.Linear(d,k)
        self.relu = nn.ReLU()
        self.out = nn.Linear(k,1)

    def forward(self,x):
        return self.out(self.relu(self.l1(x)))
```

```
[69]: model = relu_mlp(2,3)
      model.forward(X[:10,:])
```

```
[69]: tensor([[ -0.0879],
              [ 0.0092],
              [ 0.0627],
              [ 0.0424],
              [ 0.0261],
              [ 0.0063],
              [-0.1009],
              [ 0.3049],
              [ 0.0181],
              [ 0.0198]], grad_fn=<AddmmBackward0>)
```

```
[70]: Yhat = model.forward(X)
      loss = mse(Yhat,Y)
      loss
```

```
[70]: tensor(43.2361, grad_fn=<MeanBackward0>)
```

1.9 Optimization

```
[71]: def fit(model, lr, max_iter):
      optimizer = torch.optim.SGD(model.parameters(), lr=lr)
      losses=[]
      for i in range(max_iter):
          Yhat = model.forward(X)
          loss = mse(Yhat,Y)
          optimizer.zero_grad()
          loss.backward()
          optimizer.step()
          losses.append(loss.detach().numpy().item())
          if(i%10==0):
              print("%d %.2f"%(i, losses[-1]))

      return(losses)

models = {"Linear":linear(2,1),"RELU MLP":relu_mlp(2,20)}
losses = {}
for m in models:
    print("Learning model: %s"%m)
    losses[m] = fit(models[m], 0.05, 100)
    print()
```


```
Learning model: Linear
0 46.35
```

```
10 12.85
20 8.63
30 8.09
40 8.02
50 8.02
60 8.01
70 8.01
80 8.01
90 8.01
```

Learning model: RELU MLP

```
0 43.69
10 3.38
20 1.11
30 0.71
40 0.64
50 0.60
60 0.56
70 0.52
80 0.48
90 0.44
```

```
[72]: import matplotlib.pyplot as plt
plt.figure()
for m in models:
    plt.plot(losses[m])
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend(list(models.keys()))
plt.title("Training Loss vs Epoch")
plt.grid(True)
```



```
[73]: Xte = torch.randn(1000,2)
      Yte = -3*Xte[:,0] + 2*Xte[:,1]**2 + 3

      te_err={}
      for m in models:
          te_err[m]=mse(Yte,models[m].forward(Xte))
          print("%s test error: %.2f"%(m,te_err[m]))
```

Linear test error: 7.47
RELU MLP test error: 0.40